Restaurant Revenue Prediction

Odeia Busheri Simon Bruno Eddie Rudoy

Reminder: Dataset

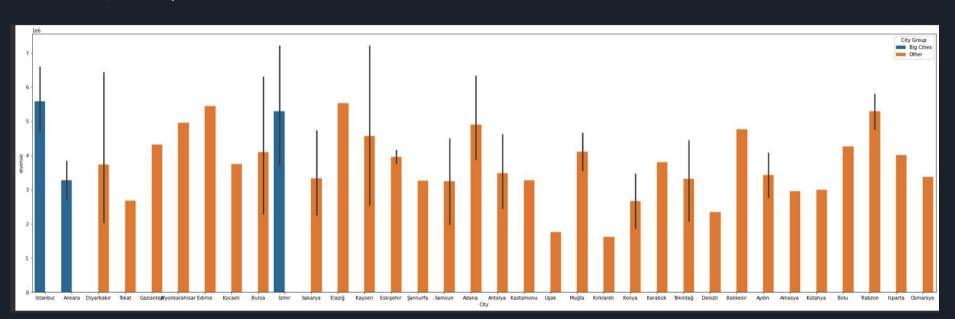
- •The train set consist of 137 rows (samples) and 43 columns which represent the features.
- Id: Restaurant id.
- •Open Date: opening date for a restaurant
- •City: City that the restaurant is in.
- •City Group: Type of the city. Big cities, or Other.
- Type: Type of the restaurant. FC: Food Court, IL: Inline, DT: Drive Thru, MB: Mobile
- •P1, P2 P37: There are three categories of these obfuscated data.

 Rank of demographic data- population in any given area, age and gender distribution, development scales.

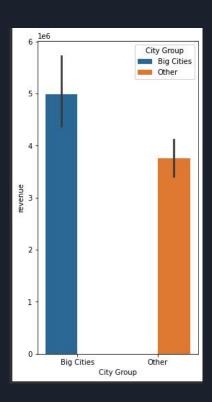
 Real estate data- mainly relate to the m2 of the location, front facade of the location, car park availability.

 Commercial data- mainly include the existence of points of interest.
- •Revenue: The revenue column indicates a (*transformed) revenue of the restaurant in a given year and is the target of predictive analysis. Transformed indicates that these values don't mean real dollar value.

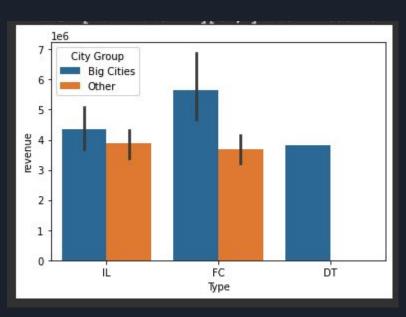
City vs Revenue



City Group vs Revenue

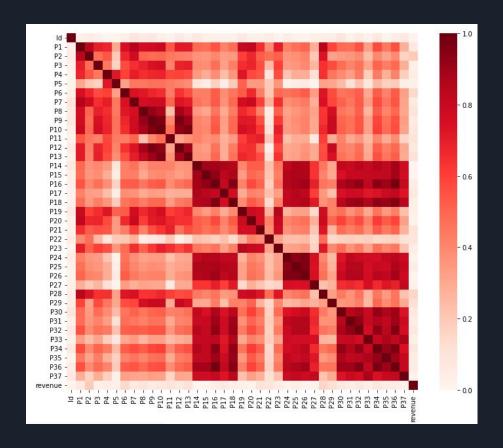


Type (City Group) vs Revenue



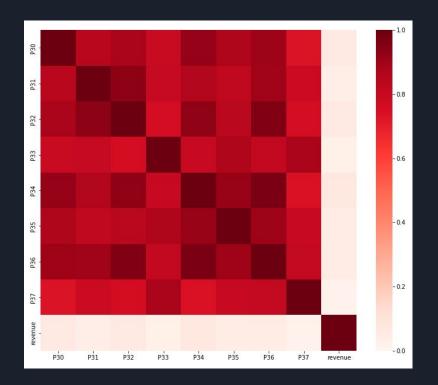
P-values heatmap

As we can see, there are groups of highly correlated features amongst each other, let's take a closer look



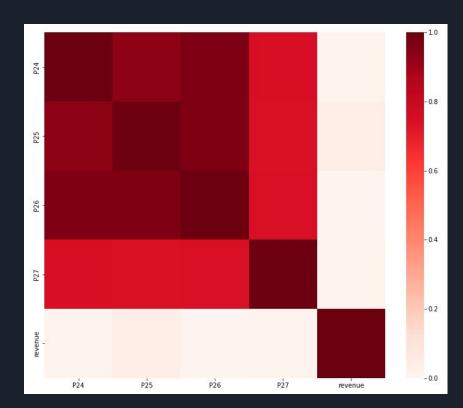
P(30-37)-values heatmap:

We can see we have a very high correlation among the P values themselves, and a low correlation to the revenue, we'll lated remove all but one



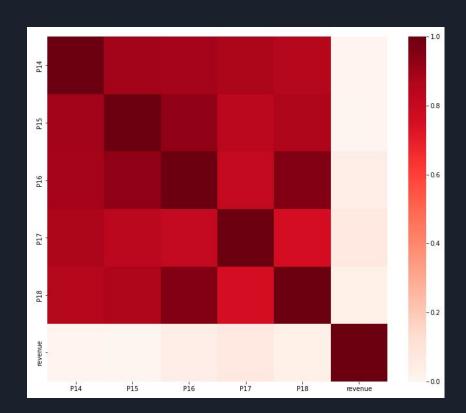
P(24-27)-values heatmap:

Same thing here



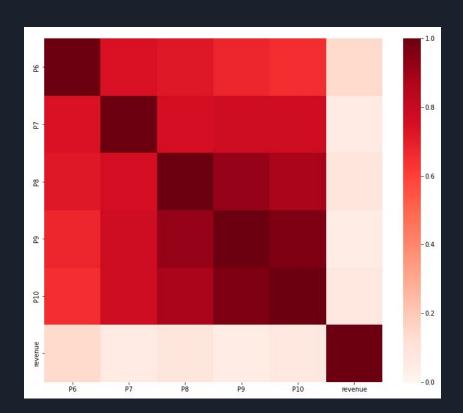
P(14-18)-values heatmap:

Also here



P(30-37)-values heatmap:

And lastly here



P variables - correlations

For each of the above groups of P-values, we'll examine the highest correlation to the revenue, and drop all but that one.

```
P14 0.006441

P15 0.000742

P16 0.037997

P17 0.067137

P18 0.034537

Name: revenue, dtype: float64

0.0671369027405342
```

```
P24 0.014222

P25 0.036365

P26 0.007650

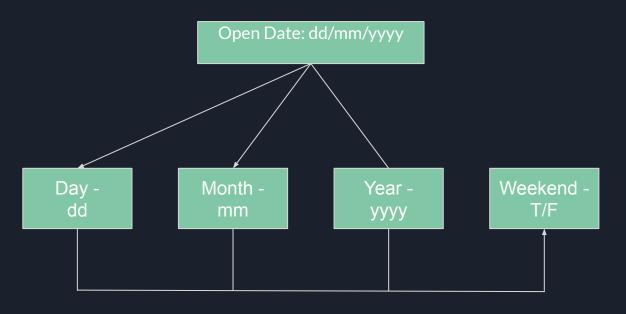
P27 0.013680

Name: revenue, dtype: float64

0.03636464638628595
```

```
P6 0.139094
P7 0.051165
P8 0.084215
P9 0.050352
P10 0.073220
Name: revenue, dtype: float64
0.13909423890893735
```

```
P30
       0.066203
P31
       0.040418
P32
       0.065857
P33
       0.032426
P34
       0.072343
       0.050156
P35
       0.050534
P36
       0.019051
P37
Name: revenue, dtype: float64
0.07234280699800917
```

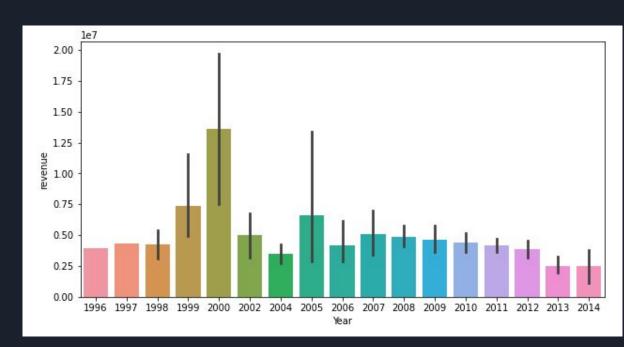


Calculate weekend from date

Let's examine how the newly extracted fields correlate to the revenue!

We'll start with the year:

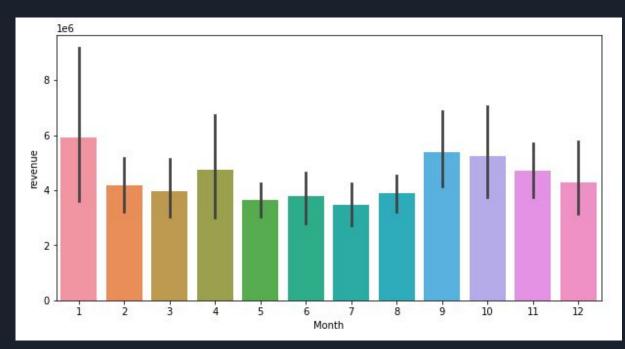
It seems that other than the start of the millenia the year didn't really correlate to the revenue



Let's examine how the newly extracted fields correlate to the revenue!

Now, let's look at the month:

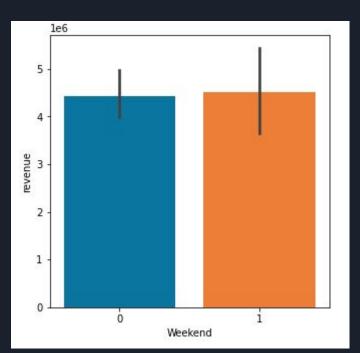
We see somewhat of a trend for colder months, but nothing that stands out significantly



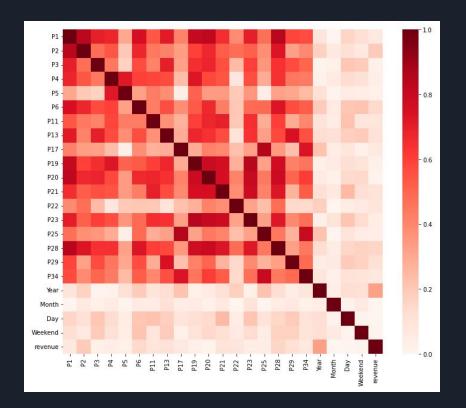
Let's examine how the newly extracted fields correlate to the revenue!

Last but not least is the weekend:

Unfortunately no large differences here as well. Some more deviation for the weekends but nothing more.



Let's take one last gander at our resulting correlation heatmap with the new date related features and the dropped P-values



Libraries & Imports

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import math
import seaborn as sns
import matplotlib.pyplot as plt
import scipy
import plotly.express as px
import plotly.graph_objects as go
# sklearn imports
import sklearn
from sklearn import metrics
from sklearn import datasets
from sklearn import pipeline
from sklearn import model_selection
from sklearn import metrics
from sklearn import pipeline
from sklearn import linear model
from sklearn import preprocessing
from sklearn import model selection
from sklearn.pipeline import Pipeline
from sklearn.pipeline import make_pipeline
from sklearn.linear_model import SGDRegressor
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import OneHotEncoder
from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.model_selection import LeavePOut
from sklearn.model_selection import RepeatedKFold
from sklearn.model_selection import KFold
from sklearn.model selection import GridSearchCV
from sklearn.model_selection import cross_val_score
```



Read Data

```
#Read data

data_path = '/content/train.csv.zip'
test_path = '/content/test.csv.zip'

restaurant_df = pd.read_csv(data_path)
test_df = pd.read_csv(test_path)

display(restaurant_df)

display(test_df)
```

Data Exploration

Check for empty values in the train set

```
print("Check if there are columns with empty values")
display(restaurant_df.isna().any())

print(f'There are {len(np.where(restaurant_df.isnull())[0])} empty values in the dataframe')

# count empty values in each column
def count_empty_values_in_each_column(df: pd.DataFrame):
    print('empty values')
    print('-----\n')

for col in df.columns:
    print(f"{col}: {df[col].isna().sum()}")

count_empty_values_in_each_column(restaurant_df)
```

empty values

There are 0 empty values in the dataframe

```
Id
              False
              False
Open Date
City
              False
City Group
              False
Type
              False
P1
              False
P2
              False
              False
P4
              False
P5
              False
P6
              False
              False
P8
              False
P9
              False
              False
              False
P12
              False
P13
              False
P14
              False
              False
P16
              False
              False
P18
              False
P19
              False
P20
              False
              False
P22
              False
              False
P24
              False
P25
              False
P26
              False
              False
P28
              False
P29
              False
P30
              False
P31
              False
P32
              False
P33
              False
P34
              False
P35
              False
P36
              False
P37
              False
revenue
              False
dtype: bool
There are 0 empty values in the dataframe
empty values
```

Check if there are columns with empty value

Divide the data to features and target

```
# divide the data to features and target
t = restaurant_df_cp['revenue'].copy()
X = restaurant_df_cp.drop(['revenue'], axis=1)
print('t')
display(t)
print()
print('X')
display(X)
```

```
5653753.0
       6923131.0
       2055379.0
       5787594.0
       9262754.0
      7217634.0
      6363241.0
Name: revenue, Length: 137, dtype: float64
137 rows x 25 columns
```

Cross Validation Score Function

```
# calculate score and loss from cv (KFold) and display graphs
def get cv score and loss(X, t, model, k=None, p=None, show score loss graphs=False, use pbar=True):
   scores losses df = pd.DataFrame(columns=['fold id', 'split', 'score', 'loss'])
   if k is not None:
       cv = KFold(n splits=k, shuffle=True, random state=1)
        raise ValueError('you need to specify k or p in order for the cv to work')
   for i, (train_ids, val_ids) in enumerate(cv.split(X)):
       X_train = X.loc[train_ids]
       t train = t.loc[train ids]
       X val = X.loc[val ids]
       t_val = t.loc[val_ids]
       model.fit(X train, t train)
       v train = model.predict(X train)
       v val = model.predict(X val)
       scores_losses_df.loc[len(scores_losses_df)] = [i, 'train', model.score(X_train, t_train), mean_squared_error(t_train, y_train, squared=False)]
       scores_losses_df.loc[len(scores_losses_df)] = [i, 'val', model.score(X_val, t_val), mean_squared_error(t_val, y_val, squared=False)]
   val scores losses df = scores losses df[scores losses df['split']=='val']
   train scores losses df = scores losses df[scores losses df['split']=='train']
   mean_val_score = val_scores_losses_df['score'].mean()
   mean_val_loss = val_scores_losses_df['loss'].mean()
   mean train score = train scores losses df['score'].mean()
   mean train loss = train scores losses df['loss'].mean()
   if show score loss graphs:
       fig = px.line(scores losses df, x='fold id', y='score', color='split', title=f'Mean Val Score: {mean val score:.2f}, Mean Train Score: {mean train score:.2f}')
       fig = px.line(scores_losses_df, x='fold_id', y='loss', color='split', title=f'Mean Val Loss: {mean_val_loss:.2f}, Mean Train Loss: {mean_train_loss:.2f}')
       fig.show()
    return mean val score, mean val loss, mean train score, mean train loss
```

Divide the cols for encoding

```
[102] numerical_cols = X.select_dtypes(include=['int64', 'float64']).columns
    categorical_cols = X.select_dtypes(include=['object', 'bool']).columns
    all_cols = np.array(X.columns)
```

```
from sklearn.compose import ColumnTransformer

ct = ColumnTransformer([
    ("encoding", OneHotEncoder(sparse=False, handle_unknown='ignore'), categorical_cols),
    ("standard", StandardScaler(), numerical_cols)])
```

One-Hot Encoding:

Transform the categorical data into few binary columns. Translate each category into a column with 0 and 1 values.

Modeling

Linear Model- SGDRegressor

XGBoost

AdaBoost

Random Forest

SGDRegressor

SGD stands for Stochastic Gradient Descent: the gradient of the loss is estimated each sample at a time and the model is updated along the way with a decreasing strength schedule (aka learning rate)

```
model_pipe = make_pipeline(ct, SGDRegressor(random_state=1))
val_score, val_loss, train_score, train_loss = get_cv_score_and_loss(X, t, model_pipe, k=10, show_score_loss_graphs=True)
print(f'mean cv val score: {val_score:.2f}\nmean cv val loss {val_loss:.2f}')
print(f'mean cv val score: {train_score:.2f}\nmean cv val loss {train_loss:.2f}')
```

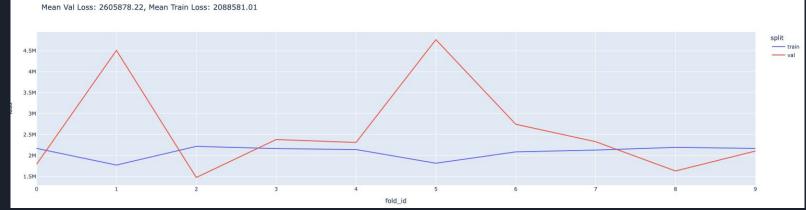
Mean cv val score: -0.56

Mean cv val loss: 2605878.22

Mean cv train score: 0.34

Mean cv val loss: 2088581.01





SGDRegressor Train Result

```
[35] SDG_reg_original = pipeline.make_pipeline(ct, SGDRegressor(random_state=1)).fit(X, t)

y_train_sdg = SDG_reg_original.predict(X)
print('Accuracy score on train', SDG_reg_original.score(X, t))
print('RMSE score on train', (mean_squared_error(t, y_train_sdg, squared=False)))

Accuracy score on train 0.31139232096021396
RMSE score on train 2129869.676142312
```

Hyper Parameter Search - using GridSearchCV

Most of our models have a lot of parameters that can be adjusted.

Each parameter value can make our model better (or worse).

We want to be able to find the best hyperparameters for our models.

Grid Search: When we want to check every parameter possible, we will use Grid Search.

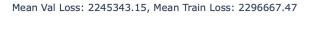
We will try all combinations of parameters and find the best one, that gives us the best score.

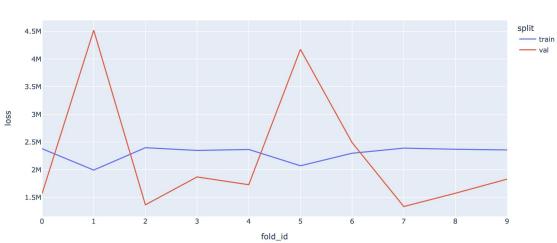
SGDRegressor Hyper Parameters

Penalty, alpha

```
[83] feature_arr = ct.fit_transform(X, t)
    feature labels = ct.get feature names out(all cols.tolist())
    X_encoded = pd.DataFrame(feature_arr, columns=feature_labels)
    hyper parameters = {'penalty': ('l2', 'l1', 'elasticnet'), 'alpha': [0.0001, 0.001, 0.01, 0.1,1]}
    gs_model = GridSearchCV(SGDRegressor(random_state=1), hyper_parameters).fit(X_encoded, t)
    print('Accuracy score for regression:')
    print('gs_model', gs_model.best_score_)
    print('best params', gs model.best params )
    Accuracy score for regression:
    gs model 0.009642228268009246
    best params {'alpha': 1, 'penalty': 'l2'}
[82] y train hyper = qs model.predict(X encoded)
    print('Accuracy score on train', gs_model.score(X_encoded, t))
    print('RMSE score on train',(mean_squared_error(t, y_train_hyper, squared=False)))
    Accuracy score on train 0.18460765025037884
    RMSE score on train 2317663.684533361
```







XGBoost - XGBRegressor

XGBoost is an optimized distributed gradient boosting library designed to be highly **efficient**, **flexible** and **portable**.

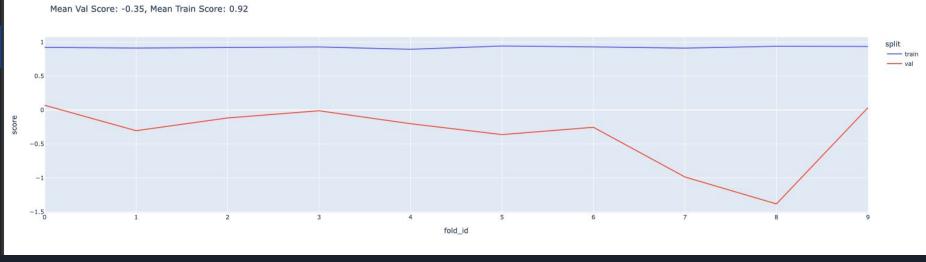
```
from xgboost import XGBRegressor
[32] XGB_model = pipeline.make_pipeline(ct, XGBRegressor(objective='reg:squarederror'))
    val_score, val_loss, train_score, train_loss = get_cv_score_and_loss(X, t, XGB_model, k=10, show_score_loss_graphs=True)
    print(f'mean cv val score: {val_score:.2f}\nmean cv val loss {val_loss:.2f}')
    print(f'mean cv val score: {train_score:.2f}\nmean cv val loss {train_loss:.2f}')
```

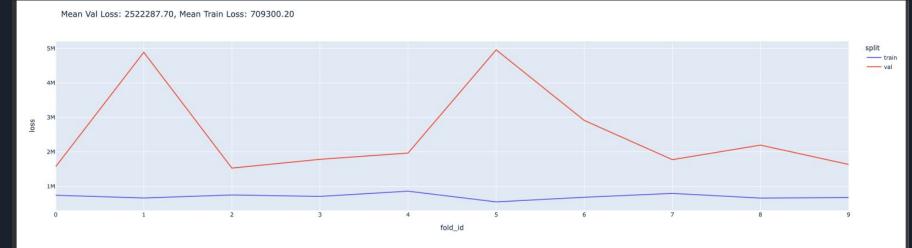
Mean cv val score: -0.35

Mean cv val loss: 2522287.70

Mean cv train score: 0.92

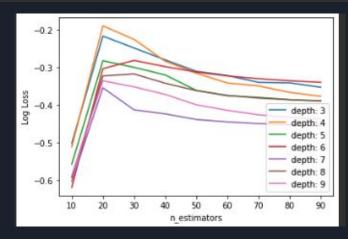
Mean cv val loss: 709300.20





XGBoost Hyperparameters

N_estimators, max_depth

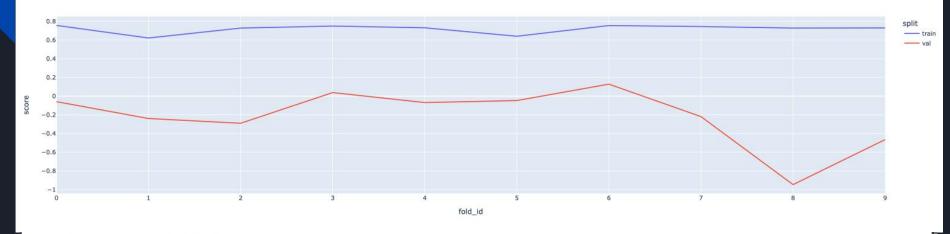


XGBoost Training Score With Hyperparameters

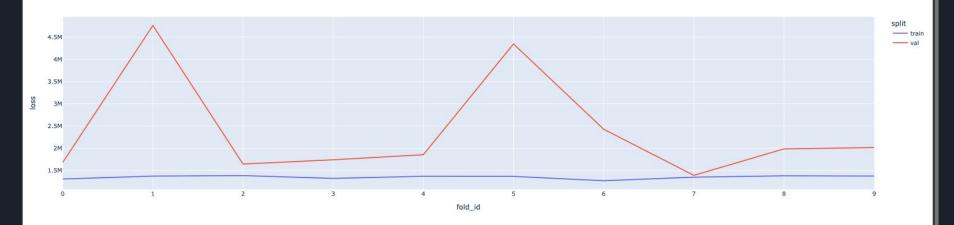
```
rs_model_pipe = pipeline.make_pipeline(ct, XGBRegressor(objective='reg:squarederror', max_depth=4, n_estimators=20)).fit(X, t)
print('Accuracy score for regression:')
y_rs_pipe = rs_model_pipe.predict(X)
print('Accuracy score on train', rs_model_pipe.score(X, t))
print('RMSE score on train', (mean_squared_error(t, y_rs_pipe, squared=False)))
# print('rs_model', rs_model_pipe.best_score_)
# print('best params', rs_model_pipe.best_params_)

Accuracy score for regression:
Accuracy score on train 0.7063737587831955
RMSE score on train 1390799.5210448876
```









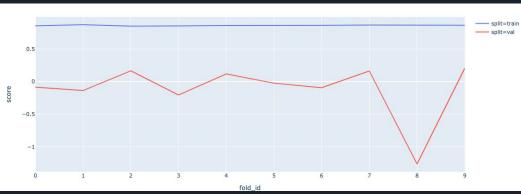
Random Forest

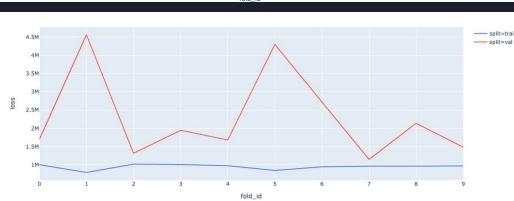
We used the basic model to get a feel for expected scores and loss values when using the Random Forest Regressor with default values using 10-fold cross validation

```
from sklearn.ensemble import RandomForestRegressor
forest_pipe = pipeline.make_pipeline(ct, RandomForestRegressor(random_state=1)).fit(X, t)

val_score, val_loss, train_score, train_loss = get_cv_score_and_loss(X, t, forest_pipe, k=10, show_score_loss_graphs=True)
print(f'mean_cv_val_score: {val_score:.2f}\nmean_cv_val_loss {val_loss:.2f}')
print(f'mean_cv_train_score: {train_score:.2f}\nmean_cv_train_loss {train_loss:.2f}')
```

Random Forest - CV graphs





mean cv val score: -0.12 mean cv val loss 2300308.81 mean cv train score: 0.86 mean cv train loss 951519.74

Random Forest - default train

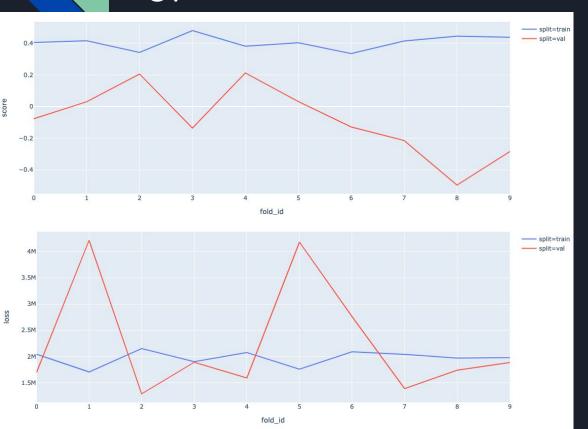
The prediction results when fitting and predicting with the default model

Accuracy score for regression: Accuracy score on train 0.8513051936335301 RMSE score on train 989726.0460099587

Random Forest - Hyperparameter selection

Choosing best hyper parameters using GridSearchCV

Random Forest - Hyperparameter selection CV



mean cv val score: -0.09 mean cv val loss 2263454.16 mean cv train score: 0.41 mean cv train loss 1971631.13

Random Forest - Hyperparameter selection Results

Chosen parameters:

```
max_depth=7
min_samples_leaf=5
min_samples_split=2
n estimators=6
```

Accuracy score for regression: Accuracy score on train 0.4505148801827643 RMSE score on train 1902589.3775937336

AdaBoost

An AdaBoost regressor is a meta-estimator that begins by fitting a regressor on the original dataset and then fits additional copies of the regressor on the same dataset but where the weights of instances are adjusted according to the error of the current prediction. As such, subsequent regressors focus more on difficult cases.

```
from sklearn.ensemble import AdaBoostRegressor

from sklearn.tree import DecisionTreeRegressor

AdaBoost_model = pipeline.make_pipeline(ct, AdaBoostRegressor(DecisionTreeRegressor(max_depth=4), n_estimators=300, random_state=1, loss='square'))

val_score, val_loss, train_score, train_loss = get_cv_score_and_loss(X, t, AdaBoost_model, k=10, show_score_loss_graphs=True)

print(f'mean cv val score: {val_score:.2f}\nmean cv val loss {val_loss:.2f}')

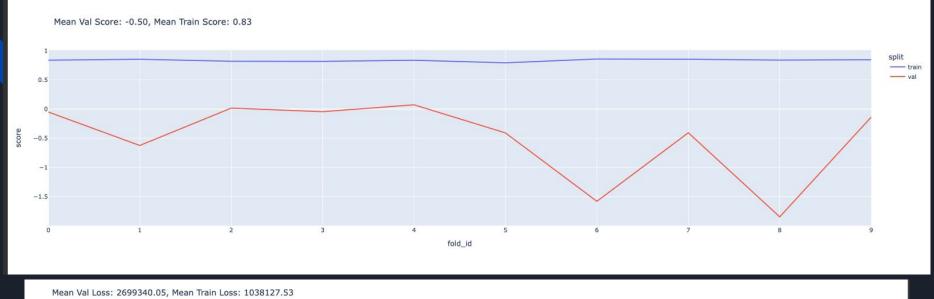
print(f'mean cv val score: {train_score:.2f}\nmean cv val loss {train_loss:.2f}')
```

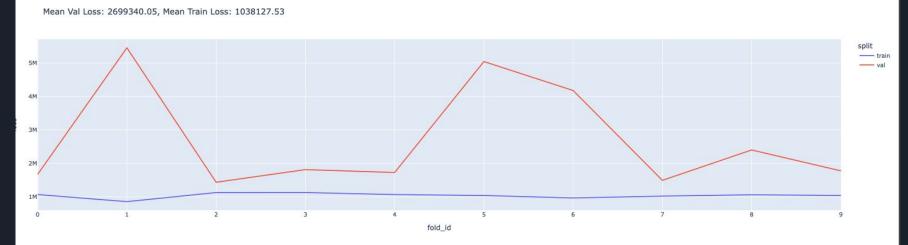
Mean cv val score: -0.50

Mean cv val loss: 2699340.05

Mean cv train score: 0.83

Mean cv val loss: 1038127.53





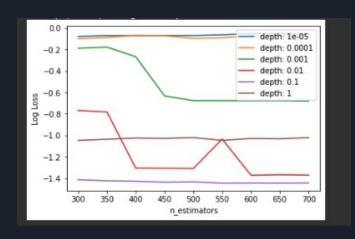
AdaBoost Train Result

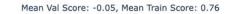
```
print('Accuracy score on train', AdaBoost_reg_original.score(X, t))
print('RMSE score on train', (mean_squared_error(t, y_train, squared=False)))
```

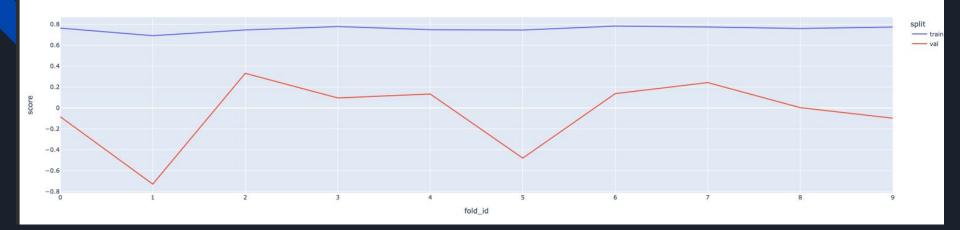
Accuracy score on train 0.8605435280405829 RMSE score on train 958487.4628856435

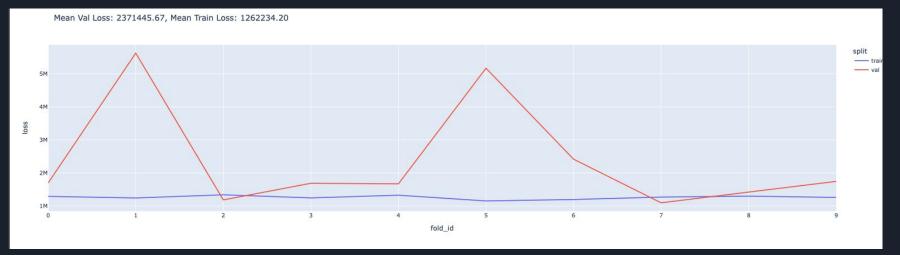
AdaBoost Hyperparameters

best params {'learning_rate': 1e-05, 'n_estimators': 600}









AdaBoost Train Hyperparameters Result

```
rs_model_pipe_Ada = pipeline.make_pipeline(ct, AdaBoostRegressor(DecisionTreeRegressor(max_depth=4), random_state=1, loss='square', learning_rate= 1e-05, n_estimators=600)).fit(X, t) print('Accuracy score for regression:')
y_rs_pipe = rs_model_pipe_Ada.predict(X)
print('Accuracy score on train', rs_model_pipe_Ada.score(X, t))
print('RMSE score on train', (mean_squared_error(t, y_rs_pipe, squared=False)))
```

Accuracy score for regression: Accuracy score on train 0.7074899962165696 RMSE score on train 1388153.4006967554

Results Comparison

	SGDRegressor	SGDRegressor tuned	XGBoost	XGBoost tuned	AdaBoost	AdaBoo st tuned	Random Forest	Random Forest tuned
R2	0.33	0.18	0.90	0.70	0.80	0.70	0.85	0.45
RMSE	2129869.67	2317663.68	785885.47	1390799.52	1130760.75	138815 3.40	989726. 04	1902589. 37

Test

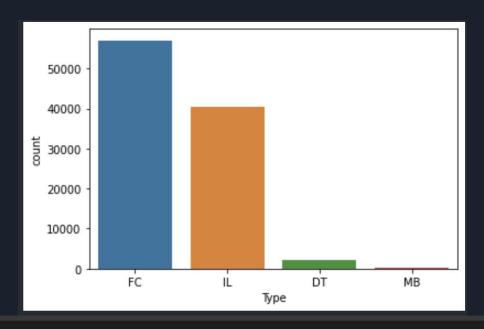
100000 samples in the test > 137 samples in the train

	ld	Open Date	City	City Group	Туре	P1	P2	РЗ	P4	P5	 P28	P29	P30	P31	P32	P33	P34	P35	P36	P37
0	0	01/22/2011	Niğde	Other	FC	1	4.0	4.0	4.0	1	2.0	3.0	0	0	0	0	0	0	0	0
1	1	03/18/2011	Konya	Other	IL	3	4.0	4.0	4.0	2	1.0	3.0	0	0	0	0	0	0	0	0
2	2	10/30/2013	Ankara	Big Cities	FC	3	4.0	4.0	4.0	2	2.0	3.0	0	0	0	0	0	0	0	0
3	3	05/06/2013	Kocaeli	Other	IL	2	4.0	4.0	4.0	2	2.0	3.0	0	4	0	0	0	0	0	0
4	4	07/31/2013	Afyonkarahisar	Other	FC	2	4.0	4.0	4.0	1	5.0	3.0	0	0	0	0	0	0	0	0
99995	99995	01/05/2000	Antalya	Other	FC	5	5.0	4.0	4.0	2	2.0	3.0	0	0	0	0	0	0	0	0
99996	99996	07/18/2011	Niğde	Other	IL	1	2.0	4.0	3.0	1	1.0	3.0	5	0	0	0	4	0	0	0
99997	99997	12/29/2012	İstanbul	Big Cities	IL	4	5.0	4.0	4.0	1	4.0	1.0	5	5	3	2	4	4	4	2
99998	99998	10/12/2013	İstanbul	Big Cities	FC	12	7.5	6.0	6.0	4	7.5	5.0	0	0	0	4	0	0	0	0
99999	99999	10/05/2010	İstanbul	Big Cities	IL	2	5.0	4.0	4.0	2	5.0	3.0	0	2	0	2	2	4	2	0

100000 rows × 42 columns

Test Data

Test set don't have empty values, how ever there is 1 for type of restaurant.



[] test_df_cp.replace('MB', 'DT', inplace=True)

Test Results Compreation

	SGDRegr essor	SGDRegressor tuned	XGBoost	XGBoost tuned	AdaBoost	AdaBoost tuned	Random Forest	Random Forest tuned
RMSE-Public	2117966.0 7873	1826302.00535	2504617.0 3774	2265356. 53008	2038560.9 43049	1911685.2 7160	2081712.88 515	1768327.26584
RMSE-Private	1928312.0 3781	1745662.51973	2340390.4 4798	2294269. 34289	2060857.1 5330	1918850.8 7706	1955265.32 349	1870204.72282